

# DSO530 Statistical Learning Methods

## Lecture 7b : Bagging, Random Forest(s) and Boosting

Dr. Xin Tong

Department of Data Sciences and Operations

Marshall School of Business

University of Southern California

[xint@marshall.usc.edu](mailto:xint@marshall.usc.edu)

# Bagging

- Deep and bushy decision trees typically suffer from high variance. And pruning can introduce some bias.
- Bootstrap aggregation, or bagging, is a general-purpose procedure for reducing the variance of a statistical learning method; we introduce it here because it is particularly useful and frequently used in the context of decision trees.
- In this approach we generate  $B$  different *bootstrapped* training data sets. We then train our method on the  $b$ th bootstrapped training set in order to get  $\hat{f}^{*b}(x)$  (not pruned), and average all the predictions:

regression  
model

bootstrapped sample  $1^* \rightarrow \hat{f}^{*1}$

Sample  $\rightarrow$  Sample  $B^* \rightarrow \hat{f}^{*B}$

side remark: let  $x_1, \dots, x_n$  be  $N(\mu, \sigma^2)$  &  $\bar{x} = \frac{x_1 + \dots + x_n}{n}$

$$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x).$$

$$\textcircled{1} x_1, \dots, x_n \text{ indep, } \text{Var}(\bar{x}) = \frac{1}{n^2} \text{Var}(x_1 + \dots + x_n) = \frac{\sigma^2}{n}$$

$$\textcircled{2} x_1 = x_2 = \dots = x_n, \text{Var}(\bar{x}) = \text{Var}(x_1) = \sigma$$

In this case averaging doesn't help at all

# Bagging

→ because they are not identical.

- Averaging these  $B$  trees reduces the “variance”. (Why?)
- $B$  is not a critical parameter with bagging; a very large value of  $B$  will not lead to overfitting.
- In practice, use  $B$  sufficiently large so that the error has settled down.
- For a given test observation in classification, we can record the class predicted by each of the  $B$  trees, and take a vote.
  - by default, the final prediction is the most commonly occurring class among the  $B$  predictions.
  - For binary classification, we can always change the threshold for  $P(Y = 1|X = x)$ ; for bagging (or random forest), this probability is computed as the fraction of  $B$  trees that predict 1.

→  $\mathbb{1}(P(Y=1|X=x) > \text{threshold})$


# Variable Importance Measures

- Although the collection of bagged trees is much more difficult to interpret than a single tree, one can obtain an overall summary of the importance of each predictor using
  - the RSS (for bagging regression trees) *try for DSO 530 project*
  - the Gini index (for bagging classification trees)
- In the case of bagging regression trees, we can record the total amount that the RSS is decreased due to splits over a given predictor, averaged over all  $B$  trees. A large value indicates an important predictor
- For bagging classification trees, we can add up the total amount that the Gini index is decreased by splits over a given predictor, averaged over all  $B$  trees.
- These variable importance measures are widely-used. However, they should be used with Caution. *Why?*

① *These measures are specific to the tree ensemble models.*

② *These measures do not take into consideration about prediction performance eval metrics*

# Random Forest(s)



- Bagging constructs trees that are too “similar” (why?), so it probably does not reduce the variance as much as we wish to.
- **Random forests** provide an improvement over bagged trees by a small tweak that *decorrelates* the trees.
- As in bagging, we build a number of decision trees on bootstrapped training samples.
- But when building these decision trees, each time a split in a tree is considered, a *random sample of  $m$  predictors* is chosen as split candidates from the full set of  $p$  predictors. The split is allowed to use *only one* of those  $m$  predictors.
- So bagging is a special case of random forest when  $m = p$ .

# Random Forest(s)

- If one does not want to spend extra efforts on  $m$ , one might use  $m = \sqrt{p}$  as a canonical choice for classification and  $m = p/3$  as a canonical choice for regression.
- As with bagging, random forests will not overfit if we increase  $B$ , so in practice people use  $B$  sufficiently large for the error rate to have settled down.
- Random forest is a really good off-the-shelf algorithm.

# Python implementation

- RandomForestClassifier and RandomForestRegressor in sklearn implement random forests in Python for classification and regression problems, respectively
- Our tutorial covers RandomForestClassifier
- Parameters:
  - `n_estimators` (default 100) is the number of trees in the forest
  - `max_features` (default `sqrt(n_features)`) is the number of features to consider when looking for the best split.  $\longrightarrow m$
- You can learn pick up RandomForestRegressor from <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html>
- RandomForestRegressor in sklearn has a default setting of `max_features=n_features`.  $\longrightarrow$  bagging ( $m=p$ )

# A definition and some questions

- If you need to communicate a one sentence ad-hoc definition of random forests:
  - *Random forests are bagged decision tree models that split on a random subset of features on each split.*
- Q: In random forest algorithms, we restrict our attention to randomly selected  $m$  out of  $p$  features in each split. Now we change this procedure to restriction to the first  $m$  features (i.e.,  $X_1, \dots, X_m$ ) in every split. Do you expect the new procedure to work well? And why?
- Q: If a decision tree partitions the feature space into regions  $R_1, \dots, R_J$ , can any of these regions be a ball?
- Q: Is random forest always a better algorithm compared to decision trees?
- Q: What are the sources of randomness that a random forest model has? Hint: 3.



# A definition and some questions

- If you need to communicate a one sentence ad-hoc definition of random forests:
  - *Random forests are bagged decision tree models that split on a random subset of features on each split.*
- Q: In random forest algorithms, we restrict our attention to randomly selected  $m$  out of  $p$  features in each split. Now we change this procedure to restriction to the first  $m$  features (i.e.,  $X_1, \dots, X_m$ ) in every split. Do you expect the new procedure to work well? And why?
- Q: If a decision tree partitions the feature space into regions  $R_1, \dots, R_J$ , can any of these regions be a ball? No
- Q: Is random forest always a better algorithm compared to decision trees? No, why? Less interpretability,
- Q: What are the sources of randomness that a random forest model has? Hint: 3. ① randomness due to sampling ② randomness due to bootstrapping

→ No, ① maybe more important features come later. (after  $m$ ) ② making similar trees makes it more like bagging than random forest. ③ randomly select  $m$  out of  $p$  features in every split. \*specific to random forest

# Boosting

- Like bagging, **boosting** is a general approach that can be applied to many statistical learning methods for regression or classification.
- Boosting is an ensemble technique where new models are added to correct the errors made by existing models.
- A differentiating characteristic Random forest: **parallel** vs. boosting: **sequential**

# A boosting algorithm for regression (optional)

---

**Algorithm 8.2** *Boosting for Regression Trees*

---

1. Set  $\hat{f}(x) = 0$  and  $r_i = y_i$  for all  $i$  in the training set.
2. For  $b = 1, 2, \dots, B$ , repeat:
  - (a) Fit a tree  $\hat{f}^b$  with  $d$  splits ( $d + 1$  terminal nodes) to the training data  $(X, r)$ .
  - (b) Update  $\hat{f}$  by adding in a shrunk version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x). \quad (8.10)$$

- (c) Update the residuals,

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i). \quad (8.11)$$

3. Output the boosted model,

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x). \quad (8.12)$$

what ???  
optional

## In practice (optional)

- XGboost
  - XGBoost is one popular implementation of boosting algorithms for its model performance.
  - It is more complicated than what we described on the previous slide. For example, subsampling and shrinkage ideas are adopted.
  - xgboost available in Python.
  - A tutorial:  
<https://xgboost.readthedocs.io/en/latest/tutorials/model.html>.
- LightGBM another popular one.
  - fast speed
  - <https://lightgbm.readthedocs.io/en/latest/Python-Intro.html>