

# Falak\_Jain\_HW2

March 20, 2022

HW2 - Falak Jain

```
[10]: import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LinearRegression
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

1.

(i) Yes

(ii) about 8-10 hours

2. Sampling balls in a box with replacement

```
[36]: np.random.seed(2)
color = ['red', 'black', 'blue', 'yellow']
for _ in range(3):
    print(np.random.choice(color, size = 5, p=[0.25, 0.25, 0.25, 0.25], replace_
↪= True))
```

```
['black' 'red' 'blue' 'black' 'black']
['black' 'red' 'blue' 'black' 'black']
['blue' 'blue' 'red' 'blue' 'red']
```

3. Mean and Variance of  $\bar{X}$

It is given that  $X_i$ 's are independent and identically distributed

Mean  $\bar{X} = (X_1 + X_2 + \dots + X_n)/n$  where  $n$  is the number of records in the dataset

$\text{Var}(\bar{X}) = \text{Var}((X_1 + X_2 + \dots + X_n)/n) / n = 1/n^2 * \text{Var}(X_1 + X_2 + \dots + X_n)$

Given that  $X_i$ 's are independent, we can say that the variance of the sum of all  $X_i = \text{sum of variances of } X_i$

$= 1/n^2 * (\text{Var}(X_1) + \text{Var}(X_2) + \dots + \text{Var}(X_n))$

$= 1/n^2 * n * \sigma^2$

$= \sigma^2 / n$

#### 4. LDA Model for Spam Detection

```
[21]: df_spam = pd.read_csv("spambase.data", header = None)
```

```
[22]: X,y = df_spam.iloc[:,0:57].values, df_spam.iloc[:,57].values

X_train, X_test, y_train, y_test = train_test_split(X,y,
                                                    test_size = 0.2,
                                                    random_state = 5,
                                                    stratify = y)
```

```
[23]: lda = LinearDiscriminantAnalysis().fit(X_train,y_train)
y_pred = lda.predict(X_test)
classification_error = np.mean(y_pred != y_test)
print(f'The classification error on the test set:␣
↪{round(classification_error*100,2)}%')
```

The classification error on the test set: 12.05%

#### 5. Plotting ROC curve and finding AUC for LDA Model

```
[30]: from sklearn.metrics import roc_curve
import matplotlib.pyplot as plt
from sklearn.metrics import auc
fpr, tpr, threshold = roc_curve(y_test, lda.predict_proba(X_test)[:,1])
roc_auc = auc(fpr, tpr)
print(f'AUC of the LDA Model: {roc_auc}')
```

AUC of the LDA Model: 0.9472782566624208

```
[27]: plt.figure()
plt.figure(figsize=(10,10))
plt.plot(fpr, tpr, color='darkorange',
lw=2, label='ROC curve (area = {0:.4f})'.format(roc_auc))
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--') # lw is linewidth
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()
```

<Figure size 432x288 with 0 Axes>

